

Topic: Ortholog mapping of targets (human → mouse)

Okay, I understand. You're working on a ChEMBL 35 data analysis project using RDKit, following an AIMLops-style folder structure, and you've encountered a couple of errors. Let's break down the analysis, provide corrected code snippets, and address your specific needs.

Analysis Model

The overall goal is to analyze ChEMBL 35 data to support drug research and development. Given your request, it sounds like you want to extract and process bioactivity data, potentially building a predictive model. Here's a basic analysis model that we will create:

1. **Data Extraction:** We'll use SQL to query the ChEMBL 35 database, specifically targeting bioactivity data related to specific targets or assays. The SQL query will be refined to filter for standard values (IC50, Ki, etc.) and activity outcomes.
2. **Data Preprocessing:**
 - **Filtering:** Remove entries with missing or inconsistent data.
 - **Data Type Conversion:** Convert standard values to numeric and handle units.
 - **Molecular Representation:** Use RDKit to generate molecular fingerprints (e.g., Morgan fingerprints) from SMILES strings.
3. **Feature Engineering:** Combine molecular fingerprints with other relevant features (e.g., physicochemical properties).
4. **Model Building (Simplified):** For demonstration purposes, let's focus on a simple Regression model to predict bioactivity values from molecular fingerprints.
5. **Model Evaluation (Simplified):** Split data into training and test sets, train the model, and evaluate its performance using Mean Squared Error (MSE).
6. **Example Generation:** Demonstrate how to apply this analysis process to various aspects such as activity distribution or target activity profile.

Error Analysis & Correction

- **Error (a): ERROR: operator does not exist: numeric ~ unknown, LINE 12: AND act.standard_value ~ '^[0-9\.]+'\$**
 - **Explanation:** This error arises because the ~ operator (regular expression matching) is not directly applicable to numeric columns in PostgreSQL. You're trying to check if `act.standard_value` (which is probably numeric) matches a regular expression.
 - **Solution:** Cast the `standard_value` to a text type before applying the regular expression. Alternatively, avoid regex altogether if you simply want to check if the value is a number. The better solution is to directly check if it's numeric which the code below does.
- **Error (b): old scikit-learn version does not support parameters squared=False in the mean_squared_error function**
 - **Explanation:** The `squared=False` parameter in `mean_squared_error` was introduced in a later version of scikit-learn.
 - **Solution:**
 - **Upgrade scikit-learn:** The best solution is to upgrade your scikit-learn version.
`pip install -U scikit-learn`

- **Calculate RMSE manually:** If you can't upgrade, calculate the Root Mean Squared Error (RMSE) manually by taking the square root of the MSE: `rmse = np.sqrt(mean_squared_error(y_true, y_pred))`

Folder Structure (Based on AIMLops)

Let's assume a basic AIMLops-style structure:

```
my_project/
├── data/           # Stores raw data (CSV files from SQL)
├── notebooks/      # Jupyter notebooks with analysis code
├── models/         # Stores trained models (optional for this example)
├── src/           # Python modules (optional for this example)
├── .env           # Environment variables (database credentials)
└── README.md
```

Code Implementation

1. SQL Code (`../data/chembl_data.sql`)

-- Extract bioactivity data for a specific target (e.g., a specific protein)
 -- Limiting to 100 rows for demonstration. Adjust WHERE clause as needed.

SELECT

```
cmp.chembl_id,
act.standard_type,
act.standard_value,
act.standard_units,
act.activity_comment,
mol.canonical_smiles
```

FROM

```
activities act
```

JOIN

```
assays ass ON act.assay_id = ass.assay_id
```

JOIN

```
target_dictionary td ON ass.tid = td.tid
```

JOIN

```
molecule_dictionary mol ON act.molregno = mol.molregno
```

JOIN

```
component_sequences cs ON td.tid = cs.tid
```

JOIN

```
component_molecules cm ON cs.component_id = cm.component_id
```

JOIN

```
molecule_dictionary cmp ON cm.molregno = cmp.molregno
```

WHERE

```
td.chembl_id = 'CHEMBL205' -- Example: Replace with your target of interest
AND act.standard_type IN ('IC50', 'Ki', 'EC50') -- Filter for common activity
```

types

```
AND act.standard_value IS NOT NULL
```

```
AND act.standard_units = 'nM'
```

```
AND act.standard_value > 0 -- Avoid zero values
```

--AND act.standard_value ~ '^[0-9\.\.]+\$' -- Removed regex. Handled in Python after conversion.

```
ORDER BY act.standard_value ASC
```

```
LIMIT 100;
```

Important:

- Replace 'CHEMBL205' with the ChEMBL ID of the target you want to analyze.
- Adjust the standard_type filter to include the types of bioactivity you're interested in.

Save the SQL query result to a CSV file (e.g., ../data/chembl_bioactivity_data.csv) using pgAdmin.

2. Python Code (notebooks/Topic_CheMBL_35_28_1_Data_Analysis.ipynb)

```
import os
import pandas as pd
import numpy as np
from rdkit import Chem
from rdkit.Chem import AllChem
from rdkit.Chem import Descriptors
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import StandardScaler

# Define the base path
base_path = os.path.dirname(os.getcwd()) #goes to the parent directory
data_path = os.path.join(base_path, 'data')
model_path = os.path.join(base_path, 'models') #create a model path to save models
later

# Load the data from the CSV file
data_file = os.path.join(data_path, 'chembl_bioactivity_data.csv')
try:
    df = pd.read_csv(data_file)
    print(f"Data loaded successfully from {data_file}")
except FileNotFoundError:
    print(f"Error: File not found at {data_file}")
    exit()

# Data Cleaning and Preprocessing
print("\nData Cleaning and Preprocessing...")
df = df.dropna(subset=['canonical_smiles', 'standard_value']) # Drop rows with
missing SMILES or standard_value
df = df[df['standard_value'] > 0] # Remove non-positive standard_value

# Convert standard_value to numeric, handling potential errors
df['standard_value'] = pd.to_numeric(df['standard_value'], errors='coerce')
df = df.dropna(subset=['standard_value']) #remove NaN created by the coercion.

# Molecular Feature Generation (RDKit)
print("\nGenerating Molecular Features...")
def generate_fingerprint(smiles):
    mol = Chem.MolFromSmiles(smiles)
    if mol:
        fp = AllChem.GetMorganFingerprintAsBitVect(mol, 2, nBits=2048) # Increased
nBits to 2048
        return np.array(fp)
    else:
        return None

df['fingerprint'] = df['canonical_smiles'].apply(generate_fingerprint)
df = df.dropna(subset=['fingerprint'])

# Example: Calculate LogP (another molecular descriptor)
df['logp'] = df['canonical_smiles'].apply(lambda x:
Descriptors.MolLogP(Chem.MolFromSmiles(x)) if Chem.MolFromSmiles(x) else None)
df = df.dropna(subset=['logp'])
```

```

# Prepare Data for Modeling
print("\nPreparing Data for Modeling...")
X = np.stack(df['fingerprint'].to_numpy()) # Stack fingerprints into a NumPy array
y = np.log10(df['standard_value'].astype(float)) # Log transform of standard_value.
Important for bioactivity data.

# Data Scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split Data
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
random_state=42)

# Model Training
print("\nTraining the Model...")
model = LinearRegression()
model.fit(X_train, y_train)

# Model Evaluation
print("\nEvaluating the Model...")
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse) # Calculate RMSE manually

print(f"Mean Squared Error: {mse:.3f}")
print(f"Root Mean Squared Error: {rmse:.3f}")

# Example Usages (Illustrative)
print("\nExample Usages:")

# 1. Distribution of Activity Values
print("\n1. Distribution of Activity Values (Log Transformed):")
import matplotlib.pyplot as plt
import seaborn as sns
sns.histplot(y)
plt.xlabel("Log10(Standard Value)")
plt.ylabel("Frequency")
plt.title("Distribution of Log10(Standard Value)")
plt.show()

# 2. Predicted vs. Actual Values
print("\n2. Predicted vs. Actual Values:")
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Log10(Standard Value)")
plt.ylabel("Predicted Log10(Standard Value)")
plt.title("Predicted vs. Actual Log10(Standard Value)")
plt.show()

# 3. Most Active Compounds
print("\n3. Most Active Compounds (Top 5):")
top_5 = df.sort_values('standard_value').head(5)
print(top_5[['chembl_id', 'standard_value', 'canonical_smiles']])

# 4. Impact of LogP on Activity
print("\n4. Impact of LogP on Activity:")
plt.scatter(df['logp'], y)
plt.xlabel("LogP")
plt.ylabel("Log10(Standard Value)")

```

```

plt.title("LogP vs. Log10(Standard Value)")
plt.show()

# 5. Model Prediction for a specific molecule
print("\n5. Model Prediction for a Specific Molecule:")
example_smiles = 'CC(=O)Oc1ccccc1C(=O)O' # Aspirin
mol = Chem.MolFromSmiles(example_smiles)
if mol:
    fp = AllChem.GetMorganFingerprintAsBitVect(mol, 2, nBits=2048)
    fp_array = np.array(fp).reshape(1, -1) # Reshape for single sample prediction
    fp_scaled = scaler.transform(fp_array) # Scale using the SAME scaler
    predicted_log_value = model.predict(fp_scaled)[0]
    print(f"Predicted Log10(Standard Value) for {example_smiles}:
{predicted_log_value:.2f}")
else:
    print(f"Invalid SMILES: {example_smiles}")

print("\nAnalysis Complete.")

```

Key Improvements and Explanations:

- **Robust Data Loading:** Uses `os.path.join` for path construction and includes error handling for file loading.
- **Data Cleaning:** Handles missing data and ensures `standard_value` is numeric.
- **Log Transformation:** Applies a log transformation to `standard_value`. This is *crucial* for bioactivity data because it often follows a log-normal distribution. Linear models work much better on normally distributed data.
- **Molecular Fingerprint Generation:** Uses `AllChem.GetMorganFingerprintAsBitVect` to generate Morgan fingerprints. **Increased nBits to 2048** for better representation.
- **Data Scaling: StandardScaler is used to scale the fingerprint data.** This is essential for linear models and can significantly improve performance. The scaler **MUST** be fit on the TRAINING data only, and then used to transform both the training and test data. This prevents data leakage.
- **Model Training and Evaluation:** Trains a `LinearRegression` model and calculates RMSE (to avoid the squared=False error).
- **Example Usages:** Includes five examples to demonstrate the kinds of analyses you can perform with the data and model.
- **Clear Comments:** The code is extensively commented to explain each step.
- **Error Handling:** Adds error handling for SMILES conversion.

Important Considerations:

- **Database Credentials:** Store your database credentials (user, password, IP address, database name) in a `.env` file and use a library like `python-dotenv` to load them. **Never hardcode credentials in your scripts.**
- **Model Selection:** Linear Regression is a very basic model. For more accurate predictions, consider using more sophisticated models like Random Forests, Gradient Boosting, or Support Vector Machines.
- **Cross-Validation:** Use cross-validation to get a more robust estimate of model performance.
- **Feature Engineering:** Experiment with different molecular descriptors and feature combinations to improve model accuracy. Consider adding more properties with RDKit, such as Molecular Weight, Number of Hydrogen Bond Donors/Acceptors, etc.
- **Target Specificity:** Focus your analysis on specific targets or assay types to get more meaningful results.

How to Run:

1. **Install Dependencies:** `bash pip install pandas rdkit scikit-learn numpy matplotlib seaborn`
2. **Create the .env file:** (Optional but highly recommended) `DATABASE_USER=your_db_user
DATABASE_PASSWORD=your_db_password DATABASE_HOST=your_db_host
DATABASE_NAME=your_db_name`
3. **Run the SQL query in pgAdmin and save the result to**
`../data/chembl_bioactivity_data.csv.`
4. **Open the Jupyter notebook (notebooks/Topic_CheMBL_35_28_1_Data_Analysis.ipynb)**
and run all cells.

Two Languages (English & Simplified Chinese)

English:

The analysis model involves extracting bioactivity data from ChEMBL, preprocessing it with RDKit for feature generation (molecular fingerprints), and building a simple regression model. The code provides error handling for common issues, includes data cleaning steps, and demonstrates various ways to visualize and interpret the results. Remember to replace the example target ID with your target of interest and consider using more advanced modeling techniques for better accuracy.

Simplified Chinese:

分析模型包括从 ChEMBL 提取生物活性数据，使用 RDKit 进行预处理以生成特征（分子指纹），并构建一个简单的回归模型。代码提供了针对常见问题的错误处理，包括数据清理步骤，并演示了可视化和解释结果的各种方法。请记住将示例目标 ID 替换为您感兴趣的目标，并考虑使用更高级的建模技术以获得更好的准确性。