

### Topic: Mapping ChEMBL drugs to WHO ATC codes

Okay, I understand. You're working on a ChEMBL 35 data analysis project using RDKit within an AIMLops framework, using a PostgreSQL database, and running your code in Jupyter notebooks. Let's break down how to approach this, address the errors, and get you running with a focused workflow.

#### 1. Analysis Model and Goal

Based on your description, it seems your primary goal is to extract and analyze data from the ChEMBL 35 database to support drug research and development. This likely involves:

- **Data Extraction (SQL):** Querying the database to retrieve relevant information about compounds, targets, activities, and other relevant data points.
- **Data Transformation and Cleaning (Python/RDKit):** Converting the data into usable formats, handling missing values, and cleaning up inconsistencies.
- **Feature Engineering (Python/RDKit):** Calculating relevant molecular properties and features using RDKit. This could include things like molecular weight, logP, topological polar surface area (TPSA), number of hydrogen bond donors/acceptors, and more.
- **Data Analysis and Modeling (Python/Scikit-learn, etc.):** Exploring relationships between features and activity, building predictive models for activity, and identifying promising compounds.
- **Visualization (Python/Matplotlib, Seaborn):** Creating plots and charts to visualize data distributions, relationships, and model performance.

**The focus on getting only 100 rows and fixing the errors suggests you're aiming for a proof-of-concept or initial exploration phase.**

#### 2. Addressing the Errors

- **Error a: ERROR: operator does not exist: numeric ~ unknown, LINE 12: AND act.standard\_value ~ '^[0-9\..]+\$'**

This error arises because you're trying to use the regular expression operator `~` to match a pattern against a numeric data type in PostgreSQL. `~` is designed for string matching. The issue is `act.standard_value` is probably defined as a numeric type (like integer or numeric).

**Solution:** Cast the `standard_value` to text before applying the regex.

- **Error b: old scikit-learn version does not support parameters `squared=False` in the `mean_squared_error` function**

This is a version compatibility issue. The `squared=False` parameter was introduced in a later version of scikit-learn.

**Solution:** There are two main approaches:

1. **Upgrade Scikit-learn:** This is the best long-term solution. In your Jupyter notebook, try: `!pip install scikit-learn --upgrade`
2. **Calculate RMSE Manually:** If upgrading isn't feasible (due to environment constraints), you can calculate the Root Mean Squared Error (RMSE) manually:

```

from sklearn.metrics import mean_squared_error
import numpy as np

```

```

mse = mean_squared_error(y_true, y_pred)
rmse = np.sqrt(mse)

```

### 3. Code Examples (SQL and Python)

Here's a structured approach with code examples. I will aim for clarity, efficiency, and addressing the errors.

```

import os
import pandas as pd
import psycopg2 # For PostgreSQL connection
from rdkit import Chem
from rdkit.Chem import Descriptors
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import numpy as np

# --- DEFINE PATHS (AIMLops Style) ---
base_path = os.getcwd() # Or your root directory
data_path = os.path.join(base_path, "data")
notebook_path = os.path.join(base_path, "notebooks") # Assuming 'notebooks' dir
model_path = os.path.join(base_path, "models")

# Ensure directories exist (or create them)
os.makedirs(data_path, exist_ok=True)
os.makedirs(notebook_path, exist_ok=True)
os.makedirs(model_path, exist_ok=True)

# --- DATABASE CONNECTION DETAILS ---
db_host = "192.168.206.136"
db_user = "rd"
db_pass = "rd"
db_name = "chembl_35"

# --- SQL QUERY ---
sql_query = """
SELECT
    act.molregno,
    act.standard_value,
    act.standard_units,
    act.pchembl_value,
    md.chembl_id,
    md.canonical_smiles
FROM
    activities act
JOIN
    molecule_dictionary md ON act.molregno = md.molregno
WHERE
    act.standard_type = 'IC50'
    AND act.standard_units = 'nM'
    AND act.standard_value IS NOT NULL
    AND act.standard_value::text ~ '^[0-9\\.]+$' -- Corrected Regex Error
LIMIT 100;
"""

# --- PYTHON CODE (Jupyter Notebook - Topic_CheMBL_35_45_1_Data_Extraction.ipynb) ---

```

```

try:
    conn = psycopg2.connect(host=db_host, user=db_user, password=db_pass,
database=db_name)
    df = pd.read_sql_query(sql_query, conn)
    conn.close()

    print("Data extracted successfully:")
    print(df.head())

    # Save to CSV
    csv_file_path = os.path.join(data_path, "chembl_ic50_data.csv")
    df.to_csv(csv_file_path, index=False)
    print(f"Data saved to: {csv_file_path}")

except Exception as e:
    print(f"Error connecting to the database or extracting data: {e}")

# --- PYTHON CODE (Jupyter Notebook -
Topic_CheMBL_35_45_2_Feature_Engineering_Modeling.ipynb) ---
# Feature Engineering and Modeling
import pandas as pd
import os
from rdkit import Chem
from rdkit.Chem import Descriptors
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import numpy as np

# --- DEFINE PATHS (AIMLops Style) ---
base_path = os.getcwd() # Or your root directory
data_path = os.path.join(base_path, "data")
notebook_path = os.path.join(base_path, "notebooks") # Assuming 'notebooks' dir
model_path = os.path.join(base_path, "models")

# --- Load the data ---
csv_file_path = os.path.join(data_path, "chembl_ic50_data.csv")
try:
    df = pd.read_csv(csv_file_path)
    print("Data loaded successfully.")
except FileNotFoundError:
    print(f"Error: File not found at {csv_file_path}")
    exit()

# --- Feature Engineering (RDKit) ---
def calculate_descriptors(smiles):
    mol = Chem.MolFromSmiles(smiles)
    if mol is None:
        return None # Handle invalid SMILES
    descriptors = {}
    descriptors["MolWt"] = Descriptors.MolWt(mol)
    descriptors["LogP"] = Descriptors.MolLogP(mol)
    descriptors["HBD"] = Descriptors.NumHDonors(mol)
    descriptors["HBA"] = Descriptors.NumHAcceptors(mol)
    descriptors["TPSA"] = Descriptors.TPSA(mol)
    return descriptors

# Apply the function to create new columns
df['descriptors'] = df['canonical_smiles'].apply(calculate_descriptors)

```

```

# Convert descriptors to separate columns
df = pd.concat([df, df['descriptors'].apply(pd.Series)], axis=1)
df.drop('descriptors', axis=1, inplace=True)

# Drop rows with NaN values in the descriptor columns (due to invalid SMILES)
df = df.dropna(subset=['MolWt', 'LogP', 'HBD', 'HBA', 'TPSA'])

# Convert IC50 to pIC50 (if pIC50 is not already present)
if 'pchembl_value' not in df.columns:
    df['pIC50'] = -np.log10(df['standard_value'] * 1e-9) # Convert nM to M

# --- Data Preparation for Modeling ---
X = df[['MolWt', 'LogP', 'HBD', 'HBA', 'TPSA']] # Feature matrix
y = df['pIC50'] # Target variable

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# --- Model Training ---
model = LinearRegression()
model.fit(X_train, y_train)

# --- Model Evaluation ---
y_pred = model.predict(X_test)

# Calculate RMSE (handling the potential scikit-learn version issue)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse) #Calculate RMSE manually

print(f"Root Mean Squared Error (RMSE): {rmse}")

# --- Save the Model (Optional) ---
import joblib
model_file_path = os.path.join(model_path, "linear_regression_model.pkl")
joblib.dump(model, model_file_path)
print(f"Model saved to: {model_file_path}")

```

### Explanation:

- **File Organization:** The code uses `os.path.join` to create paths according to your AIMLops structure, ensuring portability and reproducibility.
- **SQL Query:** The SQL query selects relevant data, filters for IC50 values in nM, and addresses the regex error by casting `standard_value` to TEXT before applying the regular expression. The `LIMIT 100` clause ensures you only retrieve 100 rows.
- **Data Loading and Cleaning:** The Python code loads the data, handles potential errors (like `FileNotFoundError`), and performs initial cleaning (dropping rows with missing descriptor values).
- **Feature Engineering (RDKit):** The `calculate_descriptors` function uses RDKit to calculate key molecular properties. Error handling is included for invalid SMILES strings.
- **pIC50 Conversion:** If `pchembl_value` is not available, the code converts IC50 values to pIC50.
- **Model Training and Evaluation:** A simple linear regression model is trained and evaluated using RMSE. The RMSE calculation is done manually to avoid the scikit-learn `squared=False` issue.
- **Model Persistence (Optional):** The trained model is saved using `joblib` for later use.

## 4. 5 Examples with Variations

Here are 5 examples demonstrating how you can modify the code:

### Example 1: Different Activity Type

- **Goal:** Extract data for Ki (inhibition constant) instead of IC50.
- **SQL Modification:**

```
SELECT
    act.molregno,
    act.standard_value,
    act.standard_units,
    act.pchembl_value,
    md.chembl_id,
    md.canonical_smiles
FROM
    activities act
JOIN
    molecule_dictionary md ON act.molregno = md.molregno
WHERE
    act.standard_type = 'Ki' -- Changed to Ki
    AND act.standard_units = 'nM'
    AND act.standard_value IS NOT NULL
    AND act.standard_value::text ~ '^[0-9\\.]+$'
LIMIT 100;
```

- **No Python code changes are needed if only data is being extracted.**

### Example 2: Different Unit

- **Goal:** Extract data for IC50 in uM (micromolar)
- **SQL Modification:**

```
sql      SELECT      act.molregno,
act.standard_value,      act.standard_units,      act.pchembl_value,
md.chembl_id,      md.canonical_smiles      FROM      activities act
JOIN      molecule_dictionary md ON act.molregno = md.molregno      WHERE
act.standard_type = 'IC50'      AND act.standard_units = 'uM' -- Changed to
uM      AND act.standard_value IS NOT NULL      AND
act.standard_value::text ~ '^[0-9\\.]+$'      LIMIT 100;
```
- **Python Modification:** The conversion to pIC50 needs to be adjusted to account for the unit change.

```
python      if 'pchembl_value' not in df.columns:      df['pIC50'] =
-np.log10(df['standard_value'] * 1e-6) # Convert uM to M
```

### Example 3: Adding More Descriptors

- **Goal:** Calculate more molecular descriptors.
- **Python Modification (in calculate\_descriptors function):**

```
def calculate_descriptors(smiles):
    mol = Chem.MolFromSmiles(smiles)
    if mol is None:
        return None # Handle invalid SMILES
    descriptors = {}
    descriptors["MolWt"] = Descriptors.MolWt(mol)
    descriptors["LogP"] = Descriptors.MolLogP(mol)
    descriptors["HBD"] = Descriptors.NumHDonors(mol)
    descriptors["HBA"] = Descriptors.NumHAcceptors(mol)
    descriptors["TPSA"] = Descriptors.TPSA(mol)
    descriptors["RotatableBonds"] = Descriptors.NumRotatableBonds(mol) #Added
```

```
descriptors["AromaticRings"] = Descriptors.NumAromaticRings(mol) #Added
return descriptors
```

You'll also need to update the X variable to include the new descriptors:

```
X = df[['MolWt', 'LogP', 'HBD', 'HBA', 'TPSA', 'RotatableBonds',
'AromaticRings']]
```

#### Example 4: Using a Different Model

- **Goal:** Use a Random Forest Regressor instead of Linear Regression.
- **Python Modification:**

```
from sklearn.ensemble import RandomForestRegressor

# --- Model Training ---
model = RandomForestRegressor(n_estimators=100, random_state=42) #Example
parameters
model.fit(X_train, y_train)
```

#### Example 5: Target Specificity by adding target information

- **Goal:** Include target name and accession in your data
  - **SQL Modification:**
- ```
sql          SELECT          act.molregno,
act.standard_value,          act.standard_units,
act.pchembl_value,          md.chembl_id,          md.canonical_smiles,
td.target_name, /* added */          td.accession /* added */          FROM
activities act          JOIN          molecule_dictionary md ON act.molregno =
md.molregno          JOIN          target_dictionary td ON act.tid = td.tid /*
added */          WHERE          act.standard_type = 'IC50'          AND
act.standard_units = 'nM'          AND act.standard_value IS NOT NULL
AND act.standard_value::text ~ '^[0-9\\.]+$'          LIMIT 100;
```

#### In summary:

I have provided a complete walkthrough of your project, from SQL data extraction to Python-based feature engineering, model building, and error handling. I've also given you 5 examples to help you adapt the code to your specific needs. Let me know if you have any more questions.

---

### Traducción al Español (Spanish Translation):

#### 1. Análisis del Modelo y Objetivo:

Basado en su descripción, parece que su objetivo principal es extraer y analizar datos de la base de datos ChEMBL 35 para respaldar la investigación y el desarrollo de fármacos. Esto probablemente implica:

- **Extracción de datos (SQL):** Consultar la base de datos para recuperar información relevante sobre compuestos, objetivos, actividades y otros puntos de datos relevantes.
- **Transformación y limpieza de datos (Python/RDKit):** Convertir los datos en formatos utilizables, manejar los valores faltantes y limpiar las inconsistencias.
- **Ingeniería de características (Python/RDKit):** Calcular propiedades y características moleculares relevantes utilizando RDKit. Esto podría incluir cosas como el peso molecular, logP, el área de superficie polar topológica (TPSA), el número de donantes/aceptores de enlaces de hidrógeno y más.

- **Análisis y modelado de datos (Python/Scikit-learn, etc.):** Explorar las relaciones entre las características y la actividad, construir modelos predictivos para la actividad e identificar compuestos prometedores.
- **Visualización (Python/Matplotlib, Seaborn):** Crear gráficos y diagramas para visualizar las distribuciones de datos, las relaciones y el rendimiento del modelo.

**El enfoque en obtener solo 100 filas y corregir los errores sugiere que está apuntando a una prueba de concepto o una fase de exploración inicial.**

## 2. Abordar los errores:

- **Error a: ERROR: operator does not exist: numeric ~ unknown, LINE 12: AND act.standard\_value ~ '^[0-9\.]+'\$**

Este error surge porque está intentando utilizar el operador de expresión regular ~ para comparar un patrón con un tipo de datos numeric en PostgreSQL. ~ está diseñado para la coincidencia de cadenas. El problema es que act.standard\_value probablemente se defina como un tipo numérico (como integer o numeric).

**Solución:** Convierta standard\_value a texto antes de aplicar la expresión regular.

- **Error b: old scikit-learn version does not support parameters squared=False in the mean\_squared\_error function**

Este es un problema de compatibilidad de versiones. El parámetro squared=False se introdujo en una versión posterior de scikit-learn.

**Solución:** Hay dos enfoques principales:

1. **Actualizar Scikit-learn:** Esta es la mejor solución a largo plazo. En su cuaderno Jupyter, pruebe: `!pip install scikit-learn --upgrade`
2. **Calcular RMSE manualmente:** Si la actualización no es factible (debido a las limitaciones del entorno), puede calcular el error cuadrático medio de la raíz (RMSE) manualmente:

```
from sklearn.metrics import mean_squared_error
import numpy as np

mse = mean_squared_error(y_true, y_pred)
rmse = np.sqrt(mse)
```

## 3. Ejemplos de código (SQL y Python)

Aquí hay un enfoque estructurado con ejemplos de código. Apunto a la claridad, la eficiencia y la resolución de los errores. (Ve el código completo arriba)

### 4. 5 Ejemplos con Variaciones

Aquí hay 5 ejemplos que demuestran cómo puede modificar el código: (Ve los ejemplos completos arriba)

En resumen, he proporcionado un recorrido completo de su proyecto, desde la extracción de datos SQL hasta la ingeniería de características basada en Python, la construcción de modelos y el manejo de errores. También le he dado 5 ejemplos para ayudarle a adaptar el código a sus necesidades específicas. Avísame si tienes más preguntas.