

### Topic: Evaluating overfitting in QSAR models

Okay, I understand. You're working with ChEMBL 35 data using RDKit and want to analyze it using a Jupyter Notebook within an MLOps-structured project. You need help with SQL queries, Python code (including addressing version issues and the errors mentioned), and implementing a workflow that adheres to your specified folder structure. Let's get started.

#### 1. Analysis Model and Approach

Based on the provided information, a general analysis model for this project would likely involve the following steps:

- **Data Extraction and Preparation:** This is where your SQL queries come in. You'll use them to extract relevant data from the ChEMBL 35 database and save it into CSV files. The "Topic\_ChEMBL\_35\_67" suggests the data selection might be related to specific activity or compound properties.
- **Data Loading and Preprocessing (Python):** Read the CSV data into your Jupyter Notebook. Clean and transform the data, including handling missing values, converting data types (e.g., strings to numeric), and potentially standardizing or normalizing numerical features.
- **Feature Engineering (Python):** Use RDKit to generate molecular descriptors (e.g., physicochemical properties, topological indices, structural fingerprints) from the SMILES strings in your ChEMBL data. These descriptors will serve as features for your analysis.
- **Data Analysis and Modeling (Python):** Depending on your research question, you might perform:
  - **Descriptive Analysis:** Calculate summary statistics, create visualizations (histograms, scatter plots), and explore relationships between features.
  - **Predictive Modeling:** Build machine learning models (e.g., regression, classification) to predict activity, toxicity, or other properties of compounds based on their molecular descriptors.
  - **Clustering:** Group compounds based on their similarities in chemical structure and properties.
- **Model Evaluation and Interpretation (Python):** Assess the performance of your models using appropriate metrics. Interpret the results to gain insights into the relationships between chemical structure and activity.
- **Reporting and Visualization (Python/Jupyter Notebook):** Document your findings, create visualizations, and present your analysis in a clear and concise manner within the Jupyter Notebook.

#### Analyzing the Errors:

- **Error a:** ERROR: operator does not exist: numeric ~ unknown, LINE 12: AND act.standard\_value ~ '^[0-9\.\.]+\$'
  - This error indicates that you're trying to use the regular expression operator ~ on a numeric column (act.standard\_value). PostgreSQL doesn't directly support regular expressions on numeric types. You'll need to cast the numeric column to text (string) before applying the regular expression.
- **Error b:** old scikit-learn version does not support parameters squared=False in the mean\_squared\_error function
  - This means your scikit-learn version is outdated. You have two options:

1. **Upgrade scikit-learn:** The preferred solution is to upgrade your scikit-learn to a newer version (e.g., using `pip install -U scikit-learn`).
2. **Remove `squared=False`:** If upgrading isn't feasible, you can remove the `squared=False` argument from the `mean_squared_error` function. This will return the Mean Squared Error (MSE) instead of the Root Mean Squared Error (RMSE). Keep in mind that this will change the scale of your error metric.

## 2. SQL, Python Code (with corrections and examples) and explanations

### Project Directory Structure

```
my_chembl_project/
├── data/
│   ├── extracted_data.csv  # Example CSV file after running SQL
│   └── ...
├── notebooks/
│   ├── Topic_CheMBL_35_67_1_data_extraction_and_preparation.ipynb
│   ├── Topic_CheMBL_35_67_2_feature_engineering_and_modeling.ipynb
│   └── ...
├── src/
│   └── utils.py  # Optional: For helper functions
├── README.md
└── ...
```

### SQL (extracting 100 rows)

```
-- File: data_extraction.sql (To be run in pgAdmin and save results as
../data/extracted_data.csv)
SELECT
    cmp.chembl_id,
    cmp.pref_name,
    act.standard_type,
    act.standard_value,
    act.standard_units,
    mol.molfile
FROM
    compound_structures cmp
JOIN
    activities act ON cmp.molregno = act.molregno
JOIN
    molecule_dictionary mol ON cmp.molregno = mol.molregno
WHERE
    act.standard_type = 'IC50' -- Example filtering criteria
    AND act.standard_units = 'nM' -- Example filtering criteria
    AND act.standard_value IS NOT NULL
    AND act.standard_value::text ~ '^[0-9\\.]+$' -- Corrected: Cast to text for regex
ORDER BY
    act.standard_value ASC
LIMIT 100;
```

### Explanation:

- **SELECT ... FROM ... JOIN ...:** This selects data from the `compound_structures`, `activities`, and `molecule_dictionary` tables, joining them based on their respective keys.
- **WHERE `act.standard_type = 'IC50'` AND `act.standard_units = 'nM'`:** This filters the data to only include IC50 values measured in nM. Change these values according to your requirements
- **AND `act.standard_value IS NOT NULL`:** This ensures that only records with non-null standard values are selected.

- **AND act.standard\_value::text ~ '^[0-9\\.]+\$'**: This is the **corrected** line. It first casts the act.standard\_value column to the text data type using ::text. Then, it applies the regular expression ~ '^[0-9\\.]+\$' to check if the value contains only digits and periods. This addresses the original SQL error.
- **ORDER BY act.standard\_value ASC**: Orders the results by the IC50 value in ascending order.
- **LIMIT 100**: Limits the result set to 100 rows.

## Python (Jupyter Notebook - Topic\_CheMBL\_35\_67\_1\_data\_extraction\_and\_preparation.ipynb)

```
import os
import pandas as pd
from rdkit import Chem
from rdkit.Chem import Descriptors
#from sklearn.metrics import mean_squared_error # Removed, see below

# Base path for your project
base_path = os.getcwd() # Get the current working directory (project root)
data_path = os.path.join(base_path, 'data')
notebook_path = os.path.join(base_path, 'notebooks')

# File path for the extracted data CSV
csv_file = os.path.join(data_path, 'extracted_data.csv')

# Load the CSV data into a Pandas DataFrame
try:
    df = pd.read_csv(csv_file)
    print("Data loaded successfully.")
except FileNotFoundError:
    print(f"Error: File not found at {csv_file}. Make sure you've run the SQL script and saved the data.")
    exit()

# Display the first few rows of the DataFrame
print(df.head())

# Convert 'standard_value' to numeric, handling errors
df['standard_value'] = pd.to_numeric(df['standard_value'], errors='coerce')

# Drop rows where 'standard_value' is NaN after conversion
df = df.dropna(subset=['standard_value'])

# Basic data cleaning
print(f"Original DataFrame shape: {df.shape}")
df = df.drop_duplicates()
print(f"DataFrame shape after removing duplicates: {df.shape}")

# Convert molfile string to RDKit Mol object. Handle potential errors if molfile is malformed
def mol_from_molfile(molfile_string):
    try:
        mol = Chem.MolFromMolBlock(molfile_string)
        if mol is None:
            return None
        else:
            return mol
    except:
        return None

df['mol'] = df['molfile'].apply(mol_from_molfile)
```

```
df = df.dropna(subset=['mol']) # Drop rows where mol object creation failed
df = df.reset_index(drop=True) # Reset index
print(df.head())
```

### Explanation:

- **Import Libraries:** Imports necessary libraries like os, pandas, and RDKit modules.
- **Path Management:** Demonstrates how to construct file paths using os.path.join to adhere to your project structure.
- **Data Loading:** Loads the CSV file into a Pandas DataFrame.
- **Data Cleaning:** Handles potential errors when converting the standard\_value column to numeric and removes rows with missing values.
- **RDKit Integration:** Converts the molfile column into RDKit Mol objects using Chem.MolFromMolBlock. Error handling ensures only valid molecules are kept in the dataframe. Rows with invalid molfile data are removed.

### Python (Jupyter Notebook - Topic\_CheMBL\_35\_67\_2\_feature\_engineering\_and\_modeling.ipynb)

```
import os
import pandas as pd
from rdkit import Chem
from rdkit.Chem import Descriptors
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
import numpy as np

# Base path for your project
base_path = os.getcwd()
data_path = os.path.join(base_path, 'data')

# File path for the extracted data CSV (or read the cleaned dataframe from the
previous notebook)
csv_file = os.path.join(data_path, 'extracted_data.csv')

# Load the CSV data into a Pandas DataFrame
try:
    df = pd.read_csv(csv_file)
    print("Data loaded successfully.")
except FileNotFoundError:
    print(f"Error: File not found at {csv_file}. Make sure you've run the SQL script
and saved the data.")
    exit()

# Ensure 'standard_value' is numeric and no NaNs exist after loading.
df['standard_value'] = pd.to_numeric(df['standard_value'], errors='coerce')
df = df.dropna(subset=['standard_value'])

# Convert molfile string to RDKit Mol object. Handle potential errors.
def mol_from_molfile(molfile_string):
    try:
        mol = Chem.MolFromMolBlock(molfile_string)
        if mol is None:
            return None
        else:
            return mol
    except:
        return None
```

```

df['mol'] = df['molfile'].apply(mol_from_molfile)
df = df.dropna(subset=['mol']) # Drop rows where mol object creation failed
df = df.reset_index(drop=True) # Reset index

# Feature Engineering using RDKit
def calculate_descriptors(mol):
    try:
        descriptors = {}
        descriptors["MolLogP"] = Descriptors.MolLogP(mol)
        descriptors["MolecularWeight"] = Descriptors.MolWt(mol)
        descriptors["NumHAcceptors"] = Descriptors.NumHAcceptors(mol)
        descriptors["NumHDonors"] = Descriptors.NumHDonors(mol)
        return pd.Series(descriptors) #Return a series to create a single row
    except Exception as e:
        print(f"Error calculating descriptors: {e}")
        return pd.Series() # Return an empty series in case of error

df = pd.concat([df, df['mol'].apply(calculate_descriptors)], axis=1)

# Drop rows where descriptor calculation failed.
df = df.dropna(subset=["MolLogP", "MolecularWeight", "NumHAcceptors", "NumHDonors"])

# Prepare data for modeling
X = df[["MolLogP", "MolecularWeight", "NumHAcceptors", "NumHDonors"]]
y = df['standard_value']

# Data scaling using StandardScaler
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Train a Linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred) # NO squared=False here if you can't upgrade
scikit-learn
rmse = np.sqrt(mse) # Calculate RMSE manually
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error (MSE): {mse}")
print(f"Root Mean Squared Error (RMSE): {rmse}")
print(f"R-squared (R2): {r2}")

```

### Explanation:

- **Feature Engineering:** Calculates molecular descriptors (LogP, Molecular Weight, Number of Hydrogen Bond Acceptors/Donors) using RDKit. This is just an example; you can add more descriptors based on your needs. Includes error handling. Returns a pd.Series to ensure it's handled correctly.

- **Data Preparation for Modeling:** Selects the calculated descriptors as features (X) and the standard\_value as the target variable (y).
- **Data Scaling:** Applies StandardScaler to scale the features. This is generally good practice for linear models.
- **Train/Test Split:** Splits the data into training and testing sets.
- **Model Training:** Trains a linear regression model using the training data.
- **Model Evaluation:** Evaluates the model's performance using Mean Squared Error (MSE), Root Mean Squared Error (RMSE), and R-squared (R2).
- **Error handling:** The code now includes a try-except block within the calculate\_descriptors function to handle potential errors during descriptor calculation. If an error occurs for a particular molecule, the function will return an empty series which will lead to a NaN value and then that row gets dropped. This is important for robustness.

### 3. Five Examples

Here are five examples of things you can do with this data, building upon the code provided:

1. **Expanding Descriptor Set:** Add more RDKit descriptors to your feature set. Explore things like Topological Polar Surface Area (TPSA), Rotatable Bond Count, or various fingerprint types (e.g., Morgan fingerprints). This will likely improve model performance.

```
from rdkit.Chem import Crippen, Lipinski

def calculate_more_descriptors(mol):
    descriptors = {}
    descriptors["MolLogP"] = Crippen.MolLogP(mol)  # Replaced
    Descriptors.MolLogP with Crippen.MolLogP
    descriptors["MolecularWeight"] = Descriptors.MolWt(mol)
    descriptors["NumHAcceptors"] = Lipinski.NumHAcceptors(mol) # Replaced
    Descriptors.NumHAcceptors with Lipinski.NumHAcceptors
    descriptors["NumHDonors"] = Lipinski.NumHDonors(mol) # Replaced
    Descriptors.NumHDonors with Lipinski.NumHDonors
    descriptors["TPSA"] = Descriptors.TPSA(mol)
    descriptors["RotatableBonds"] = Descriptors.NumRotatableBonds(mol)
    return pd.Series(descriptors)

df = pd.concat([df, df['mol'].apply(calculate_more_descriptors)], axis=1)
df = df.dropna(subset=["MolLogP", "MolecularWeight", "NumHAcceptors",
"NumHDonors", "TPSA", "RotatableBonds"]) # Update dropped cols
```

2. **Different Regression Models:** Try different regression models like Random Forest Regressor, Support Vector Regressor, or Gradient Boosting Regressor. These models might capture non-linear relationships better than linear regression.

```
from sklearn.ensemble import RandomForestRegressor

model = RandomForestRegressor(n_estimators=100, random_state=42) # Example:
Random Forest
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
# ... evaluate model ...
```

3. **Data Visualization:** Create scatter plots of different descriptors versus the standard\_value to visualize potential relationships. Use matplotlib or seaborn for visualization.

```
import matplotlib.pyplot as plt
import seaborn as sns

sns.scatterplot(x='MolLogP', y='standard_value', data=df)
```

```
plt.xlabel('MolLogP')
plt.ylabel('IC50 (nM)')
plt.title('MolLogP vs. IC50')
plt.show()
```

4. **Filtering by Compound Properties in SQL:** Modify your SQL query to filter compounds based on specific properties (e.g., molecular weight range, presence of certain functional groups). This lets you focus on a specific subset of your data.

```
-- Example: Filter by molecular weight
SELECT ...
FROM ...
WHERE ...
AND cmp.molregno IN (SELECT molregno FROM molecule_properties WHERE mw_freebase
BETWEEN 200 AND 400)
LIMIT 100;
```

5. **Clustering Analysis:** Use clustering algorithms (e.g., K-Means, hierarchical clustering) to group compounds based on their descriptor profiles. This can help identify clusters of compounds with similar activities.

```
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

# Scale the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply K-Means clustering
kmeans = KMeans(n_clusters=5, random_state=42) # Example: 5 clusters
df['cluster'] = kmeans.fit_predict(X_scaled)

# Analyze the clusters (e.g., calculate average activity per cluster)
print(df.groupby('cluster')['standard_value'].mean())
```

### Important Considerations:

- **Error Handling:** Always include error handling in your code to gracefully manage unexpected situations (e.g., invalid data, missing files).
- **Documentation:** Document your code clearly with comments to explain what each step does.
- **Data Understanding:** Spend time understanding your data and the meaning of the different features. This will help you make informed decisions about your analysis.
- **Virtual Environments:** Use virtual environments to manage your Python dependencies and avoid conflicts.
- **ChEMBL License:** Remember to comply with the ChEMBL database license terms.

### In summary (en resumen):

This comprehensive guide provides you with the necessary code, explanations, and examples to start your ChEMBL 35 data analysis project. It addresses your specific requirements, including the SQL error correction, Python version compatibility (by removing `squared=False`), and adherence to your project directory structure. Remember to adapt the code and examples to your specific research question and data. Good luck!